
CHURN IS DEAD

The Supervisor / Closer Split Audit

Map your CSM week into supervisor, closer, and babysitter buckets — then score the judgment ratio that exposes whether you redesigned the job or just bolted agents onto a dying one.

QB didn't kill the CSM role. It split it into two jobs almost nobody was hired to do — and one trap most teams sleepwalk into. This audit forces you to classify every task in a CSM's week and confront the number that tells you which half you're building.

Scoring: Rate each item 1-5. 1 = not happening, 5 = embedded in operations.

Created by: Kuber Sethi

Source: churnisdead.com

Version: 1.0 -- June 2026

SECTION 1

Curation Work — The Supervisor

The agent drafts 40 check-ins; three are wrong in ways only a human who knows the account can catch. The supervisor is the reason the other 37 don't erode trust — but only if they can catch them before send.

Criteria	What Good Looks Like	What Bad Looks Like
Pre-send review on risk accounts	Agent output to named-risk accounts hits an approval gate; a human judges it before the customer sees it	Human gets a notification after the email has already shipped and read
Account-context judgment	Supervisor catches the message congratulating a laid-off champion or referencing a deprecated feature	Errors only surface when the customer replies angrily or churns
Correction feedback loop	When a supervisor overrides a draft, the correction trains the next run	Overrides vanish into the void; the agent repeats the same mistake next week
Coverage of high-stakes moments	Renewal asks, post-Sev-1 outreach, and exec-touch messages are always human-curated	The agent autonomously sends a renewal ask the week after a Sev-1 outage
Supervisor capacity is staffed for	Headcount and time are explicitly allocated to curation as real work	Curation is assumed to happen 'in the cracks' of an already-full generalist week

Score This Section

Criteria	Score
Approval gates exist on anything touching a named-risk account	1-5
Humans catch account-specific errors before send, not after	1-5
Corrections train the agent's next run instead of disappearing	1-5

Your Score: ____ / 5

SECTION 2

Commercial Work — The Closer

No agent navigates a procurement freeze, a champion change, a security review, or a multi-stakeholder budget fight. The bottleneck was never detecting the problem — it's the human org on the customer's side, and software can't navigate a CAB.

Criteria	What Good Looks Like	What Bad Looks Like
Procurement & budget navigation	A human owns the strategy on frozen renewals and board-level budget fights	The agent emails stakeholders at midnight while the renewal sits frozen for months
Champion change response	When the champion leaves, a closer rebuilds the relationship map and re-establishes value	Departure is noted in the CRM; outreach continues to an empty seat
Security & compliance blockers	A human drives the data-residency or security review that froze the expansion	The blocker is flagged six weeks ago and nobody is navigating the customer's CAB
Multi-stakeholder politics	Closer reads internal knife-fights over platform ownership and picks a side strategically	Internal politics are invisible to the role; the account is treated as one monolith
Escalation judgment	Closer makes the call to escalate to the customer's CIO or put a human on a plane	Escalation only happens after the account is already in churn motion

Score This Section

Criteria	Score
A named human owns commercial strategy on frozen/at-risk renewals	1-5
Champion changes trigger active relationship rebuilding, not just a CRM update	1-5
The role navigates the customer's internal org, not just sends messages to it	1-5

Your Score: ____ / 5

SECTION 3

The Babysitter Trap

Supervising agent output with no authority to change anything isn't supervision — it's busywork you're now liable for. If the human enters the loop after the action completes, you didn't buy supervision; you bought liability laundering.

Criteria	What Good Looks Like	What Bad Looks Like
Halt power	Human can stop a send before it goes out, not apologise for it after	Human can only acknowledge, escalate, or add a note to an action already taken
Rewrite power	Human can change the cadence, segment logic, or message when they spot a pattern	Human files a ticket with ops and waits; the workflow keeps running unchanged
Authority matches accountability	The human is accountable only for output they could actually prevent	Senior-CSM money is paid for someone to feel responsible for a machine's decisions
Vendor loop placement	The 'human in the loop' enters before the action completes	The human is a reviewer of completed actions hard-coded as a babysitter
Time spent watching vs. deciding	Little of the week is spent reading sent emails or watching dashboards refresh	Hours go to reviewing decisions the machine already made and shipped

Score This Section

Criteria	Score
The human can halt a pending action, not just react to a completed one	1-5
The human can rewrite cadence/logic directly without a ticket queue	1-5
Accountability is scoped to actions the human had authority to prevent	1-5

Your Score: ____ / 5

SECTION 4

Role Redesign & The Judgment Ratio

The judgment ratio is the share of the week spent making decisions an agent cannot make versus babysitting ones it already made. Above 60% you've redesigned the job. Below 40% you built a babysitter and you're paying senior money for it.

Criteria	What Good Looks Like	What Bad Looks Like
Explicit supervisor vs. closer staffing	Roles are designed and hired for supervisor or closer work distinctly	One generalist is expected to do both while pretending nothing changed
Judgment ratio is measured	Each task is classified supervisor/closer/babysitter and the ratio is tracked	Activity is measured by outreach volume — messages sent, not problems prevented
Ratio above 60%	The human does only what a human can; the agent carries the administrative rest	Below 40% — the role has no judgment in it and is a babysitter with a title
Proactive = prevented, not sent	Success is measured by problems prevented and renewals navigated	Green dashboard counts messages that left the outbox; nobody acted on the insight
Role survival logic	Leadership can articulate which half of the split each role is built for	Leadership calls it 'strategic transformation' while bolting agents onto a dead job

Score This Section

Criteria	Score
Tasks are explicitly bucketed and the judgment ratio is calculated	1-5
The role's judgment ratio sits above 60%	1-5
Hiring and role design distinguish supervisor from closer	1-5

Your Score: ____ / 5

CHURN IS DEAD

Your Scorecard

Dimension	Your Score	Max
1. Curation Work — The Supervisor	_____	5
2. Commercial Work — The Closer	_____	5
3. The Babysitter Trap	_____	5
4. Role Redesign & The Judgment Ratio	_____	5
TOTAL	_____	20

Total Score	Assessment	What It Means
16-20	Uncuttable	Embedded, measurable, strategically irreplaceable.
12-15	Defensible	Solid but gaps remain. Focus on lowest dimension.
8-11	Vulnerable	Real risk ahead. Treat as urgent.
4-7	Critical	Cannot survive scrutiny. Major changes needed.

What to do next: Take your lowest-scoring section and build a 30-day action plan. One at a time. Then share this with leadership -- not to ask permission, but to show you know where the gaps are.

If your CSM job could be done by an agent emailing 83 people at midnight, it was already a dead job. The agent didn't kill it. It made it visible.

-- Kuber

churnisdead.com | Weekly frameworks that replace hope with strategy.